

Libgcc: When exceptions collide

Max Neunhöffer



A Rabbit Hole

This is the story of a rather deep rabbit hole...

- ▶ assembler debugging (endless loops),
- ▶ linked lists (always good for trouble),
- ▶ mutexes and concurrency,
- ▶ Linux perf tools,
- ▶ C++ exceptions and their runtime support,
- ▶ libmusl, static linking,
- ▶ weak references in ELF, ...

The problem

We use **continuous integration** with **Jenkins**.

A **PR must only** be merged, if Jenkins is green.

In my branch, I had **mysterious** shutdown problems.

More details about the testing

Jenkins runs:

- ▶ Lots of integration tests (distributed system)
- ▶ very high load (many concurrently running tests)
- ▶ Timeout would produce **gigabytes of core files**

Furthermore, we use:

- ▶ **Static binaries built with libmusl**

Program terminated with signal SIGABRT, Aborted.

```
#0  0x00007f4d1442cc57 in __deregister_frame_info_bases (begin=0x7f4d14adaf60)
    at /home/buildozer/aports/main/gcc/src/gcc-8.3.0/libgcc/unwind-dw2-fde.c:222
222 /home/buildozer/aports/main/gcc/src/gcc-8.3.0/libgcc/unwind-dw2-fde.c: No
such file or directory.
```

(gdb) bt

```
#0  0x00007f4d1442cc57 in __deregister_frame_info_bases (begin=0x7f4d14adaf60)
    at /home/buildozer/aports/main/gcc/src/gcc-8.3.0/libgcc/unwind-dw2-fde.c:222
#1  __deregister_frame_info_bases (begin=0x7f4d14adaf60)
    at /home/buildozer/aports/main/gcc/src/gcc-8.3.0/libgcc/unwind-dw2-fde.c:201
#2  0x00007f4d1199b6ce in __do_global_ctors_aux ()
#3  0x0000000000000027 in ?? ()
#4  0x00007f4d1199e192 in runServer (argc=0, argv=0x0, context=...)
    at /work/ArangoDB/arangod/RestServer/arangod.cpp:284
#5  0x00007ffddc78f898 in ?? ()
#6  0x00007f4d14447af5 in _fini ()
#7  0x0000000000000000 in ?? ()
```

The main function has already terminated!

Detective work in a core

By default, no source code for this is available.

This is **annoying**, one does not want to rebuild `libgcc` with debug symbols.

This meant **assembler debugging**.

0x7f4d1442cc46 <__deregister+103>	lea 0x28(%rbx),%rax
0x7f4d1442cc4a <__deregister+107>	mov 0x28(%rbx),%rbx
0x7f4d1442cc4e <__deregister+111>	test %rbx,%rbx
0x7f4d1442cc51 <__deregister+114>	je 0x7f4d1442cc82 <__deregister+163>
0x7f4d1442cc53 <__deregister+116>	mov 0x18(%rbx),%rdi
0x7f4d1442cc57 <__deregister+120>	testb \$0x1,0x20(%rbx)
0x7f4d1442cc5b <__deregister+124>	jne 0x7f4d1442cc6b <__deregister+140>
0x7f4d1442cc5d <__deregister+126>	cmp %rdi,%rbp
0x7f4d1442cc60 <__deregister+129>	jne 0x7f4d1442cc46 <__deregister+103>

%rip was 0x7f4d1442cc46,
 %rbx was 0x7f4d1568e460 and x/6xg 0x7f4d1568e460 showed:

0x7f4d1568e460 :	0x00007f4d11995000	0x0000000000000000
0x7f4d1568e470 <object.6149+16>:	0x0000000000000000	0x00007f4cd6f00b00
0x7f4d1568e480 <object.6149+32>:	0x000000002190d0d9	0x00007f4d1568e460

We are looking at an endless loop!!!

The **libgcc** source code

Is here:

<https://github.com/gcc-mirror/gcc/blob/4ac50a4913ed81cc83a8baf865e49a2c62a5fe5d/libgcc/unwind-dw2-fde.c#L221>


```
struct object
{
    void *pc_begin;
    void *tbase;
    void *dbase;
    union {
        const struct dwarf_fde *single;
        struct dwarf_fde **array;
        struct fde_vector *sort;
    } u;

    union {
        struct {
            unsigned long sorted : 1;
            unsigned long from_array : 1;
            unsigned long mixed_encoding : 1;
            unsigned long encoding : 8;
            unsigned long count : 21;
        } b;
        size_t i;
    } s;

    struct object *next;
};
```

Proving the code to be correct

Being a mathematician, the next I did was:

Prove the code to be correct.

Surely, I would find the bug in `libgcc...`

Catching the bug in the act

Monitoring showed (after the fact):

one CPU was spinning at 100%

How could I catch it?

Linux perf tools to the rescue:

```
sudo perf probe -x build/bin/arangod -a "unwind=unwind-dw2-fde.c:1072 ob p"  
sudo perf record -e "probe_arangod*" -aR --call-graph dwarf  
    (hit Control-C after the recording)  
sudo perf script > perf.history
```

Gotcha! ... Really???

After **many hours** of continuous running...

(using an identical machine and identical testing load)

It happened again!

MaintenanceWork 8985 [025] 4713776.942241: probe_arangod:unwind: (7f4d1442cd77) ob=0x7f4d1568e460
p=0x7f4d165b4590

```
2c93d77 _Unwind_Find_registered_FDE (inlined)
2c93d77 _Unwind_Find_FDE (/work/ArangoDB/build/bin/arangod)
2c91474 uw_frame_state_for (/work/ArangoDB/build/bin/arangod)
2c92065 uw_init_context_1 (/work/ArangoDB/build/bin/arangod)
2c924a9 _Unwind_RaiseException (/work/ArangoDB/build/bin/arangod)
2c3ea25 __cxa_throw (/work/ArangoDB/build/bin/arangod)
520c33 replicationSynchronize (/work/ArangoDB/build/bin/arangod)
5251d3 arangodb::maintenance::SynchronizeShard::first (/work/ArangoDB/build/bin/arangod)
```

MaintenanceWork 8981 [012] 4713776.992994: probe_arangod:unwind: (7f4d1442cd77) ob=0x7f4d1568e460
p=0x7f4d1568e488



```
2c93d77 _Unwind_Find_registered_FDE (inlined)
2c93d77 _Unwind_Find_FDE (/work/ArangoDB/build/bin/arangod)
2c91474 uw_frame_state_for (/work/ArangoDB/build/bin/arangod)
2c92065 uw_init_context_1 (/work/ArangoDB/build/bin/arangod)
2c924a9 _Unwind_RaiseException (/work/ArangoDB/build/bin/arangod)
2c3ea25 __cxa_throw (/work/ArangoDB/build/bin/arangod)
520c33 replicationSynchronize (/work/ArangoDB/build/bin/arangod)
5251d3 arangodb::maintenance::SynchronizeShard::first (/work/ArangoDB/build/bin/arangod)
```

Impossible - protected traffic by mutex

What must have happened:

- ▶ The same exception was thrown **for the first time** at **approximately the same time** (50 ms difference).
- ▶ **The same object** was inserted into the linked list **twice**.

A → B → C → nullptr



Mutex - or no mutex - that is the question

BUT: There is a Mutex! Or is there?

<https://github.com/gcc-mirror/gcc/blob/4ac50a4913ed81cc83a8baf865e49a2c62a5fe5d/libgcc/unwind-dw2-fde.c#L1046>


```
static inline int  
__gthread_mutex_lock (__gthread_mutex_t *__mutex)  
{  
    if (__gthread_active_p ())  
        return __gthrw_(pthread_mutex_lock) (__mutex);  
    else  
        return 0;  
}
```

<https://github.com/gcc-mirror/gcc/blob/4ac50a4913ed81cc83a8baf865e49a2c62a5fe5d/libgcc/gthr-posix.h#L744>

```
#ifdef __GLIBC__
__gthrw2(__gthrw__(__pthread_key_create),
        __pthread_key_create,
        pthread_key_create)
# define GTHR_ACTIVE_PROXY __gthrw__(__pthread_key_create)
#elif defined (__BIONIC__)
# define GTHR_ACTIVE_PROXY __gthrw__(pthread_create)
#else
# define GTHR_ACTIVE_PROXY __gthrw__(pthread_cancel)
#endif
```

```
static inline int
__gthread_active_p (void)
{
    static void *const __gthread_active_ptr
        = __extension__ (void *) & GTHR_ACTIVE_PROXY;
    return __gthread_active_ptr != 0;
}
```

The crucial finding...

The music plays here:

<https://github.com/gcc-mirror/gcc/blob/4ac50a4913ed81cc83a8baf865e49a2c62a5fe5d/libgcc/gthr-posix.h#L156>

and here:

<https://github.com/gcc-mirror/gcc/blob/4ac50a4913ed81cc83a8baf865e49a2c62a5fe5d/libgcc/gthr-posix.h#L89>

Linking is black magic of the worst kind

A **weak reference** to sth in an object file is one

- ▶ that is resolved to something, **if that sth is present**
- ▶ and to **nullptr** otherwise.

Here, on Linux, a weak ref to `pthread_create` is used.

Rationale:

- ▶ A multi-threaded program links `pthread_create`,
- ▶ A single-threaded program **does not**.

When exactly does it happen?

- ▶ We are on **Linux**,
- ▶ we are **not** using **Glibc**,
- ▶ we are not on Android with the Bionic C-library,
- ▶ we **statically link** the executable,
- ▶ we do not explicitly use `pthread_cancel`,
- ▶ the **first exception** that happens, happens **in two threads exactly at the same time**
- ▶ in one thread the execution is suspended at an unlucky time during the list manipulation in `_Unwind_Find_FDE`.

Whose fault is it?

This is an **unholy alliance** between

- ▶ over-optimization for the single-threaded case,
- ▶ subtle differences between **GLibC** and **libmusl**,
- ▶ using black linking magic and the preprocessor and
- ▶ exception handling complexity.

Alles wird gut...

I reported this problem:

- ▶ to **libgcc** first, who said it is a bug in **libmusl**,
- ▶ then to **libmusl**, who said it is a bug in **libgcc**,
- ▶ then an argument broke out and it turned out that a corresponding fix was done previously for **libgfortran** and **libstdc++**
- ▶ finally fixed for gcc-10 in **libgcc**.

Lessons learned

- ▶ **Do not over-optimize** at the cost of unnecessary complexity.
- ▶ C++ runtime and **exception handling is complicated**.
- ▶ C++ exceptions can be quite expensive when actually thrown
- ▶ **C-code** can be **hard to understand**, in particular in the presence of lots of platform-specific special cases and **preprocessor magic**.
- ▶ The **Linux perf tools** are an incredible **asset** for debugging.
- ▶ **Open source** makes such an investigation possible in the first place.
- ▶ The teams developing these fundamental tools like **gcc** and **libmusl** are **very responsive** and actually help when problems arise.

Thank you!

Resources

- ▶ Blog: <https://www.arangodb.com/2019/09/when-exceptions-collide/>
- ▶ Gcc bug report: https://gcc.gnu.org/bugzilla/show_bug.cgi?id=91737
- ▶ Libmusl bug report: <https://www.openwall.com/lists/musl/2019/09/17/1>